Alex Morgan

AI Engineer | LLM Systems, RAG, Computer Vision

+1 555 010 0200 | alex.morgan@example.com | City, Country

Summary

AI Engineer with 3+ years of experience delivering LLM and computer-vision systems from research to production. Builds RAG, multi-agent, voice, OCR, and edge-AI products end to end, from model design and evaluation to reliable services and deployment. Add one or two differentiating achievements, publications, or domain specialties here.

Experience

Example AI Company

2024 – Present

AI Engineer

- Built **Project Atlas**, a multi-agent assistant for a domain workflow; designed intent routing, retrieval, and response orchestration with an LLM application framework and Python services.
- Delivered a document-ingestion and RAG pipeline that indexes unstructured knowledge, combines semantic retrieval with structured data, and supports multi-turn user questions.
- Productionized model-serving workflows with asynchronous APIs, a relational database, and observability; replace this sentence with your measurable latency, quality, cost, or adoption result.

Example Product Studio

2022 – 2024

Machine Learning Engineer

- Designed and deployed conversational AI and voice-assistant experiences for customer-facing workflows, integrating LLMs, speech services, and external APIs.
- Implemented hybrid retrieval that combines vector search and metadata filtering to improve access to enterprise knowledge.
- Engineered backend services for concurrent sessions and real-time processing; state the scale or performance target when available.

Example Research Lab

2021 – 2022

AI Researcher

- Developed a lightweight vision model for on-device classification; replace with task, F1 or accuracy, throughput, model size, and compute budget.
- Evaluated latency-aware architectures for edge deployment and documented reproducible experiments and ablations.
- Optimized inference with TensorRT and TFLite for resource-constrained devices.

Publications

Example Paper Title: A concise description of the research contribution
Venue or journal | First author | DOI or URL

Optional Second Publication or Thesis Title
Venue, thesis, or work in progress | Author role | DOI or URL

Technical Skills

LLM & Agents
RAG, GraphRAG, multi-agent systems, LLM frameworks, commercial and open-source models

ML & Vision
Python, PyTorch, TensorFlow, scikit-learn, OpenCV, OCR, edge AI

Retrieval & Data
Vector database, graph database, SQL, NoSQL, object storage

Production & MLOps
FastAPI, Docker, model serving, ONNX, orchestration, CI/CD, Linux, Git

Education

Example University
City, Country

B.Eng. in Computer Science | Specialization or research focus
2017 – 2021 | GPA: 3.70/4.00

Certificates

Relevant Certification or Language Test
Score or credential ID

Issuing Organization

Optional Second Certification
Score or credential ID

Issuing Organization

Honors & Awards

Example Research or Engineering Award
Award level

Awarding organization
Year